

Highlights

A physics-informed digital twin for tube-life-aware operation of industrial steam methane reformers under uncertainty

Izbassar Orynbassar, Madina Sissenbay

- Open digital twin couples reformer kinetics, furnace, tube wall and creep life
- Calibrated on published plant data: outlet T within 4 K, wall T within 12 K
- Firing, steam-to-carbon and excess air govern hot-spot temperature and life
- Catalyst ageing costs 2.3 times more tube life than any flexible-load scenario
- Creep-curve scatter dominates life uncertainty; regime rankings survive it

A physics-informed digital twin for tube-life-aware operation of industrial steam methane reformers under uncertainty

Izbassar Orynbassar^{a,*}, Madina Sissenbay^b

^a*School of Information Studies, Syracuse University, 343 Hinds Hall, Syracuse, 13244, NY, USA*

^b*Institute of Petrochemical Engineering and Ecology named after N.K. Nadirov, Safi Utebayev Atyrau University of Oil and Gas, Atyrau, Kazakhstan*

Abstract

Background. Catalyst tubes are the life-limiting component of a steam methane reformer, yet operating decisions are rarely taken with a quantified tube-life cost. **Methods.** A one-dimensional Xu–Froment tube model, a long-furnace model of a top-fired box furnace, a Larson–Miller creep-life module and Gaussian-process surrogates with conformal prediction intervals are coupled into a digital twin built entirely from published data and open code. The tube model is verified against the industrial case of Xu and Froment, the coupled model is calibrated on three published plant operating points with leave-one-case-out cross-validation, and the twin is exercised through global sensitivity analysis, hourly annual scenarios, multi-objective optimisation and Monte Carlo uncertainty propagation. **Results.** Calibration reproduces outlet temperatures within 4 K and tube-wall temperatures within 12 K. Firing intensity, steam-to-carbon ratio and excess air explain 82 % of the hot-spot temperature variance with negligible interactions. Under outlet-temperature control, load-following and renewable-following cost no tube life per kmol of hydrogen, whereas catalyst ageing under a constant methane-slip target multiplies the life cost by 2.3 within four years. A Pareto knee delivers +16 % hydrogen at –18 % fuel per kmol and 0.28 of the base life cost, and creep-curve scatter contributes 62 % of the life-prediction variance. **Conclusions.** The twin ranks operating regimes robustly, but absolute lives require alloy-specific rupture data, and the sign of the life cost of flexible operation is decided by the flame-length response to load, which is identified as the key measurement need.

Keywords: steam methane reforming, digital twin, creep life, conformal prediction, multi-objective optimisation, uncertainty quantification

1. Introduction

Steam methane reforming (SMR) supplies the large majority of the world’s hydrogen, close to 100 Mt per year and still produced predominantly from natural gas without carbon capture [1]; it will remain the dominant route for at least a decade while electrolytic capacity is built up. The core of an SMR plant is a top-fired or side-fired furnace in which several hundred centrifugally cast high-alloy tubes, filled with nickel catalyst, carry natural gas and steam at 25–35 bar while their outer walls are heated to 1100–1200 K. The tubes operate in the creep regime by design, and their remaining life is the single most important maintenance variable of the plant: a tube failure forces an unplanned shutdown, retubing costs about a tenth of the installed plant cost [2], and service life varies between 30 000 and 180 000 h with material quality and operating conditions [3]. Tube life is governed almost entirely by the outer-wall temperature at the hottest elevation, because the rupture time roughly halves for every 10–20 K of additional metal temperature [2, 4]; the placeholder curve used here halves the life every 11 K at the stress of the calibrated tube.

*Corresponding author.

Email addresses: iorynbas@syr.edu (Izbassar Orynbassar), m.sissenbay@aogu.edu.kz (Madina Sissenbay)

URL: <https://orcid.org/0009-0002-4224-3249> (Izbassar Orynbassar), <https://orcid.org/0009-0009-8160-7472> (Madina Sissenbay)

Operating decisions therefore trade hydrogen production, fuel consumption and tube life against each other, but the trade-off is rarely quantified: plants are run against a maximum tube-wall temperature limit read from pyrometer surveys, and the effect of a load, steam-to-carbon or burner change on tube life is left to experience. Two developments make a quantitative approach both necessary and possible. Hydrogen plants are increasingly expected to follow variable downstream demand or to operate alongside intermittent electrolytic hydrogen, which raises the question of what tube life flexible operation costs; and validated models of the tube [2, 5–7] and furnace [8, 9] exist, surrogates make them fast enough for hourly simulation and optimisation [10–13], and rupture data for the relevant alloys are documented in public data sheets and manufacturers’ design curves [14, 15, 24]. What is missing is a coupled, calibrated and uncertainty-aware chain from operating inputs to tube life, built entirely on published data and open code, so that its conclusions can be checked and its data replaced.

1.1. Previous work

The reformer tube has been modelled at every level of complexity. Xu and Froment [5] established the Langmuir–Hinshelwood–Hougen–Watson kinetics of methane steam reforming and water-gas shift on a $\text{Ni/MgAl}_2\text{O}_4$ catalyst that remains the standard intrinsic model, and in the companion paper [6] simulated an industrial tube with a one-dimensional heterogeneous model. The tube has since been coupled to a zone model of the furnace [8], to a unified top-fired model with three-dimensional zonal analysis [9] and to a dynamic heterogeneous formulation [7]. Latham and co-workers [2, 17] calibrated a Hottel zone-method model of a top-fired box reformer on plant measurements including infrared tube-wall temperature surveys, and published the operating points and absolute tube-wall temperatures in the thesis appendix [17]; that data set is the basis of the present calibration. Computational-fluid-dynamics models of the furnace [18–20] resolve the temperature field but are used to balance tube temperatures rather than to quantify life, and the data-driven work of the last decade — furnace-balancing and smart-manufacturing frameworks [10, 19, 20], CFD-trained machine-learning surrogates [11, 13] and uncertainty-aware production optimisation [12] — addresses production and efficiency: tube life, when it appears, enters as a constraint on a maximum temperature rather than as an objective with a quantified cost, and the uncertainty in kinetics, heat transfer and creep data is not propagated to the conclusions.

On the materials side, the creep behaviour of the 25Cr–20Ni (HK40) and 25Cr–35Ni–Nb (HP-Nb) families of centrifugally cast tube alloys is characterised in the long-term data sheets of the National Institute for Materials Science [14, 15, 21] and in remaining-life studies of service-exposed tubes [3, 22, 23]. Yeh [24] combined a computed tube-wall temperature with a manufacturer’s Larson–Miller curve for a proprietary HP-Nb alloy to estimate life, and showed that evaluating the life at the average instead of the maximum tube temperature overstates it by two orders of magnitude; the same curve is used here as a placeholder. The Larson–Miller parameter [25] and the linear life-fraction rule of Robinson [26] remain the basis of design codes such as API 530 [4].

1.2. Contribution and research questions

This paper presents a physics-informed digital twin of an industrial top-fired SMR that links operating inputs to tube creep life through a one-dimensional tube model, a long-furnace model, a Larson–Miller life module, machine-learning surrogates with distribution-free prediction intervals, and Monte Carlo uncertainty propagation (Fig. 1). The tube model is verified against the industrial case of Xu and Froment [6], the coupled model is calibrated on the plant data of Latham [17] with cross-validation, and every number in the results is produced by a versioned open pipeline from the published data. The twin is then used to answer four research questions:

- RQ1** Which operating variables govern the hot-spot temperature and the creep-life consumption of the tube, and how do they interact?
- RQ2** What is the tube-life cost, per unit of hydrogen produced, of flexible operation (load-following, renewable-following, over-firing events) and of catalyst ageing, relative to steady operation?
- RQ3** Which operating regimes are Pareto-optimal in hydrogen production, fuel consumption and life consumption, and how does the answer depend on how steam is charged?
- RQ4** How uncertain are these answers, which uncertainty sources dominate, and which conclusions survive the uncertainty?




Figure 1: (a) Cross-section of a top-fired box reformer: burners in the roof between rows of catalyst tubes, flames releasing heat in the upper part of the firebox, flue gas flowing downward co-currently with the process gas and leaving through tunnels at the floor. (b) Module chain of the digital twin: operating inputs are mapped through the furnace and tube models, which exchange the outer-wall temperature profile and the heat flux, to the creep-life module; surrogates and uncertainty quantification wrap the chain.

Section 2 describes the models, the data and the numerical methods. Section 3 reports verification, calibration and the answers to RQ1–RQ4. Section 4 draws the industrial implications and states the limitations, and Section 5 concludes.

2. Methods

The twin is implemented as a Python package (`rdt`) whose modules follow the chain of Fig. 1b. All physical data (kinetic parameters, plant operating points, creep curves) are read from annotated YAML and CSV files; no number is hard-coded in the model code. Thermodynamic properties, heat capacities and reaction equilibrium constants are taken from Cantera [27] with the GRI-Mech 3.0 species data, so that kinetics and energy balance are mutually consistent.

2.1. Reactor tube model

The tube is described by a steady one-dimensional pseudo-homogeneous plug-flow model with axial coordinate z from the inlet (top) to the outlet. The state consists of the molar flows F_i of CH_4 , H_2O , H_2 , CO , CO_2 and N_2 , the process-gas temperature T and the pressure p . The intrinsic kinetics are those of Xu and Froment [5] for steam reforming (r_1), water-gas shift (r_2) and overall reforming to CO_2 (r_3):

$$r_1 = \frac{k_1}{p_{\text{H}_2}^{2.5}} \frac{p_{\text{CH}_4} p_{\text{H}_2\text{O}} - p_{\text{H}_2}^3 p_{\text{CO}} / K_1}{\text{DEN}^2}, \quad (1)$$

$$r_2 = \frac{k_2}{p_{\text{H}_2}} \frac{p_{\text{CO}} p_{\text{H}_2\text{O}} - p_{\text{H}_2} p_{\text{CO}_2} / K_2}{\text{DEN}^2}, \quad (2)$$

$$r_3 = \frac{k_3}{p_{\text{H}_2}^{3.5}} \frac{p_{\text{CH}_4} p_{\text{H}_2\text{O}}^2 - p_{\text{H}_2}^4 p_{\text{CO}_2} / K_3}{\text{DEN}^2}, \quad (3)$$

$$\text{DEN} = 1 + K_{\text{CO}} p_{\text{CO}} + K_{\text{H}_2} p_{\text{H}_2} + K_{\text{CH}_4} p_{\text{CH}_4} + K_{\text{H}_2\text{O}} p_{\text{H}_2\text{O}} / p_{\text{H}_2}, \quad (4)$$

with partial pressures in bar. The rate coefficients k_j and adsorption constants K_i follow the Arrhenius and van't Hoff forms with the activation energies, adsorption enthalpies and reference values of Table 5 of [5], reparameterised at the reference temperatures 648 K and 823 K used in that table so that the reported confidence intervals can be propagated directly (Section 2.7). The equilibrium constants K_1 – K_3 are computed from Gibbs energies of formation rather than from fitted expressions, which removes a source of inconsistency with the energy balance (Section S1).

The balances per tube are

$$\frac{dF_i}{dz} = A_{cs} \rho_b \sum_j \nu_{ij} \eta_j r_j, \quad (5)$$

$$\sum_i F_i c_{p,i} \frac{dT}{dz} = \pi d_i U (T_{wo}(z) - T) - A_{cs} \rho_b \sum_j \eta_j r_j \Delta H_j(T), \quad (6)$$

$$-\frac{dp}{dz} = \frac{1-\phi}{\phi^3} \left(1.75 + 150 \frac{1-\phi}{Re_p} \right) \frac{\rho_g v_s^2}{d_p}, \quad (7)$$

where A_{cs} is the cross-section, ρ_b the bed density, ν_{ij} the stoichiometric coefficients, η_j the effectiveness factors, ϕ the bed voidage, d_p the equivalent particle diameter, v_s the superficial velocity and Re_p the particle Reynolds number; Eq. (7) is the Ergun equation [28]. Rather than solving the pellet problem, constant effectiveness factors $\eta_j = 0.05$ in the first tenth of the tube and 0.1 thereafter are used, and higher alkanes in the feed (C_2 – C_5) are converted to CO and CH_4 at the inlet; both follow Latham et al. [2] and are set out in Section S1.

The overall heat-transfer coefficient referred to the inner surface combines the bed-side film coefficient α_i and the conduction through the tube wall,

$$\frac{1}{U} = \frac{1}{\alpha_i} + \frac{d_i}{2\lambda_t} \ln \frac{d_o}{d_i}, \quad T_{wi} = T_{wo} - q''_o \frac{d_o}{2\lambda_t} \ln \frac{d_o}{d_i}, \quad (8)$$

with the tube conductivity $\lambda_t = 29.6 \text{ W m}^{-1} \text{ K}^{-1}$ of [17]. The film coefficient uses the Leva–Grummer correlation [30] in the form written by Latham et al. [2], multiplied by a heat-transfer factor f_{htg} that is estimated from plant data (Section 2.5). The alternative chain of Xu and Froment [6], based on De Wasch and Froment [31] and Yagi and Kunii [32], is implemented as an option and discussed in Section 3.1. Equations (5)–(7) are integrated with the LSODA solver of SciPy for a prescribed outer-wall profile $T_{wo}(z)$ (stand-alone mode) or as part of the coupled system below.

2.2. Furnace model and coupling

The furnace is a top-fired box with rows of burners in the roof between rows of tubes (Fig. 1a; the reference geometry has 7 rows of 48 tubes and 8 rows of 12 burners [2]). Instead of the zone method of Hottel and Sarofim [33] used by Plehiers and Froment [8] and Latham et al. [2], we use a one-dimensional long-furnace model in the sense of Roesler [34]: the flue gas is a plug flow travelling downward co-currently with the process gas, and the gas–tube radiative exchange is lumped into one effective exchange factor F_{gt} . Per unit height z (zero at the roof),

$$\dot{n}_{fg} c_{p,fg} \frac{dT_{fg}}{dz} = r(z) - q(z) \pi d_o N_t, \quad (9)$$

$$q(z) = \sigma F_{gt} (T_{fg}^4 - T_{wo}^4) + h_c (T_{fg} - T_{wo}), \quad (10)$$

$$q(z) = U (T_{wo} - T) \frac{d_i}{d_o}, \quad (11)$$

where N_t is the number of tubes, q the flux per unit outer tube area and h_c a furnace-side convective coefficient, which at furnace-gas temperatures is a minor correction to the radiative flux. The heat-release density $r(z)$ is a downward-opening parabola that vanishes at $z = L_q L$, with the dimensionless flame length L_q and the top-zone release fraction α_{top} as calibration parameters. Because a pure plug flow starting from the mixed fuel–air inlet temperature would give an unphysical cold region before the first heat release, the top zone is treated as well mixed: the fraction α_{top} of the heat is released instantly at the roof. The release profile, the convective coefficient and the combustion calculation are given in Section S1.

Equation (11) closes the system: at every z the outer-wall temperature T_{wo} is the root of the flux balance between Eqs. (10) and (11), so that furnace and tube are integrated together in one downward sweep (Section S1). The global energy closure of the coupled solution is verified to 10^{-5} of the combustion heat, and one coupled operating point takes about 0.3 s on a laptop.

2.3. Tube wall and creep life

The creep-life module converts the outer-wall temperature profile into a life estimate. The hoop stress is computed from the mean-diameter formula of API 530 [4], $\sigma_h = p(d_o - t)/(2t)$ with wall thickness t , at the process-gas pressure of each elevation; the thin-wall and Lamé variants differ from it by less than 12 % at the geometry used ($d_o = 146$ mm, $t = 15$ mm, 28 bar: 10.8, 12.2 and 12.4 MPa respectively). The wall thickness is not given by Latham; it is taken as 15 mm, consistent with the 15.3 mm wall of the industrial tube of Xu and Froment [6], and varied between 12 and 18 mm in the uncertainty analysis. The temperature entering the life calculation is the outer-wall temperature, which is conservative with respect to the mid-wall temperature; the difference is shown in Fig. 4.

The rupture time follows the Larson–Miller parameter [25],

$$\text{LMP} = 10^{-3} T (C + \log_{10} t_r), \quad \log_{10} \sigma = P_3(\text{LMP}), \quad (12)$$

with T in K, t_r in h, $C = 22.96$ and a cubic master curve P_3 fitted to the minimum-property rupture points of the manufacturer’s curve for Manaurite XM reproduced by Yeh [24]. This curve is a *placeholder*: the alloy of the reference plant is not stated, and a manufacturer’s design curve carries no scatter information. The module is written so that the curve can be replaced by a fit to any rupture data set that carries scatter information, such as manufacturers’ published average and minimum design curves for HP-Nb tube alloys or rupture data tabulated in the open literature (Section 2.4). Damage under a piecewise-constant history is accumulated with the linear life-fraction rule of Robinson [26], $D = \sum_k \Delta t_k / t_r(T_k, \sigma_k)$, and rupture is predicted at $D = 1$. Throughout the paper, life is reported as a consumption rate $1/t_r$ at the hot spot, or as damage per kmol of hydrogen, relative to the calibrated base case, because ratios of rupture times depend on the slope of the master curve at the operating point rather than on its absolute position and are therefore far less sensitive to the placeholder than absolute values are.

2.4. Data

Verification data. The industrial simulation of Xu and Froment [6] (Fig. 3 of that paper) provides axial profiles of methane and CO₂ conversion, process-gas, inner-wall and outer-wall temperatures and pressure for one operating point with the feed of their Table 2 (520 °C, 29 bar, steam-to-carbon 3.36, tube 0.1016 m inner and 0.1322 m outer diameter, 11.12 m heated length). The curves were digitised with WebPlotDigitizer with 120–260 points per curve and separate axis calibrations for conversion, pressure and temperature; the calibration resolution is about 2 K and 0.005 in conversion. The data are stored as tidy CSV files with their metadata (Section S4 of the Supplementary material). The outer-wall curve serves as the boundary condition and the remaining curves as targets.

Plant data. Latham [17] reports, in Appendix H, four operating points of an industrial top-fired reformer (Plants A, B and two states of Plant C) with inlet conditions, feed and fuel compositions, outlet temperature and composition, flue-gas outlet temperature and the tube-wall temperatures measured by infrared pyrometer at two peep-hole elevations (3.66 and 8.53 m from the top of the 12.5 m heated length). Table 1 lists the cases. The instrument uncertainties stated in the thesis, 2 K for the outlet temperature, 3 K for the wall temperatures and 8 K for the flue gas, are treated as one-sigma weights in the calibration; the standard deviation of the wall temperature across the tube population at a given elevation is much larger, 15–22 K in the journal version of the data [2], and is used to interpret the residuals. Plant C1 is excluded from calibration because its feed flow is inconsistent with the other cases (it equals the reported combustion-air flow to four digits and exceeds the Plant A feed by 22 % at a lower plant rate), which we attribute to a transcription error in the source. The excess air of each case is derived from the reported fuel, off-gas and air flows and is 21.6 % for Plant A.

Creep data. The placeholder Larson–Miller curve is transcribed from the minimum-property design curve for Manaurite XM reproduced by Yeh [24]. An ingestion routine is provided that fits the Larson–Miller constant C and the master curve by cross-validation with per-heat offsets and estimates the heat-to-heat scatter from the residuals; it is validated on a synthetic data set with known C and scatter, both recovered within 10 %, and it accepts manufacturers’ average and minimum design curves without heat identifiers, in which case the average-to-minimum gap is used as the scatter band. The long-term rupture data of the National Institute for Materials Science for the HK40 and HP40 alloys [14, 15] are the natural reference for this module, but they are served for interactive viewing only and their terms of use prohibit the assembly of consolidated data sets for machine-learning purposes [16]; the present work therefore relies on published design curves and on rupture data tabulated in the open literature, and only derived quantities (curve coefficients, C , scatter) are stored in the repository.

2.5. Calibration and cross-validation

Four parameters of the coupled model cannot be taken from the literature: the exchange factor F_{gt} , the flame length L_q , the top-zone release fraction α_{top} and the heat-transfer factor f_{htg} . They are estimated by weighted non-linear least squares on the three plant cases against nine residuals per case, each scaled by its measurement uncertainty; the residual set, the optimiser, the multi-start procedure and the covariance estimate are given in Section S1. Predictive performance is assessed by leave-one-case-out cross-validation: the parameters are refitted on two cases and the held-out case is predicted. The four-parameter fit places α_{top} at its bound with a correlation of 0.985 to L_q , so the adopted configuration fixes $\alpha_{top} = 0.182$, the value fitted by Latham et al. [2] to the same furnace with the zone method, and estimates the remaining three; the increase in χ^2 is reported in Section 3.1.

2.6. Surrogates and conformal prediction

The coupled model takes about 0.3 s per operating point, which is adequate for calibration and sensitivity analysis but not for hourly annual histories, evolutionary optimisation or Monte Carlo over histories. Surrogates are trained on a Latin-hypercube design of 2000 points over the seven operating inputs (Table S1) for eight outputs, and the model with the lowest five-fold cross-validated error is selected per target and evaluated on an independent Saltelli design of 4096 points. The targets, the three model families compared, their hyperparameters and the two physics-derived features are given in Section S3. Prediction intervals are obtained by the cross-conformal (CV+) method [37] as implemented in MAPIE [38], which provides finite-sample marginal coverage without a distributional assumption; the Gaussian-process native intervals are reported for comparison.

2.7. Sensitivity, scenarios, optimisation and uncertainty propagation

RQ1. Global sensitivity uses the variance-based Sobol indices [39] estimated with SALib [40] on the Saltelli design evaluated with the physics model, together with an open-loop Latin-hypercube campaign and a closed-loop grid in steam-to-carbon ratio and load in which the firing is solved so that the outlet temperature equals the Plant A value (1105.55 K).

RQ2. Annual histories of 8760 hourly quasi-steady states are evaluated through the surrogates, with every 24th hour and every hour outside the training bounds re-evaluated by the physics model. Two control modes are represented by inverting the surrogates for the firing factor with a vectorised bisection: holding the outlet temperature at its base value, or holding the dry methane slip. Six scenarios cover steady operation, daily and renewable-following load variation, over-firing events, four years of catalyst ageing under each control mode, and two steam-to-carbon settings; their definitions are given in Table S2. Damage is accumulated hourly by Eq. (12) and Robinson’s rule, and each scenario is also evaluated at its annual-mean inputs to quantify the error of a steady-state assessment. Start-up and shutdown transients, thermal fatigue and creep–fatigue interaction are outside the scope.

RQ3. The non-dominated sorting genetic algorithm NSGA-II [41] in pymoo [42] maximises net hydrogen per tube and minimises firebox fuel per kmol H_2 and life consumption per kmol H_2 over five decision variables (load, steam-to-carbon ratio, inlet temperature, excess air, firing) subject to a methane slip not above the base value, a hot-spot temperature not above 1193 K and an outlet temperature not above 1125 K, using the Gaussian-process surrogates (population 200, 300 generations); the basis of the 1193 K limit is given in Section S2. Every Pareto member is re-evaluated with the physics model and the constraints are re-checked. Because firebox fuel per kmol H_2 ignores the cost of raising the additional steam, a steam credit α is introduced: the total fuel counts firebox fuel plus a fraction $(1 - \alpha)$ of the heat needed to raise the process steam and to preheat the dry feed (enthalpy basis in Section S2), so that $\alpha = 1$ recovers the firebox-only objective and $\alpha = 0$ charges steam in full. The iso-production front (hydrogen within 1 % of the base) is recomputed for α from 0 to 1.

RQ4. Uncertain inputs are grouped into kinetics (the three rate parameters and four adsorption constants of Xu and Froment with normal distributions from their reported 95 % intervals, a log-uniform factor of 0.5–2 on the effectiveness factors and ± 0.03 on the catalyst activity), heat transfer (the calibrated F_{gt} , L_q and f_{htg} with their joint covariance), creep (a normal scatter of 0.3 decades on $\log_{10} t_r$ representing a factor of two between heats, C uniform within ± 1 of 22.96 with the master curve refitted for each C , wall thickness 12–18 mm and tube conductivity ± 15 %) and measurement (a uniform ± 10 K offset on the pyrometer-referenced hot-spot temperature). A scrambled Sobol




Figure 2: Verification against the digitised industrial case of Xu and Froment [6] (symbols) with the digitised outer-wall temperature as boundary condition: (a) methane and CO_2 conversion, (b) process-gas and inner-wall temperatures, (c) pressure and (d) temperature residuals, for the matched film coefficient $\alpha_i = 348 \text{ W m}^{-2} \text{ K}^{-1}$ (solid lines) and for the Leva-Grummer correlation with $f_{\text{htg}} = 1$ (dotted).

design of 2048 samples is propagated through the physics model at eight operating points taken from RQ2 and RQ3; the same sample is used at every point so that ratios between points are paired. Variance is decomposed by varying one group at a time (1024 samples per group) and, independently, by Sobol indices computed on a Gaussian-process surrogate of the joint sample. Finally, the fraction of paired samples for which each qualitative conclusion of RQ2 and RQ3 holds is reported, both with the calibrated flame length and with a bounding structural variant in which the flame length shrinks with load as $L_q \text{ load}^{0.5}$; the variant is a bound rather than an expectation, for the reasons set out in Section S2. The complete pipeline from design to figures runs in 35 minutes on a laptop.

3. Results

3.1. Verification and validation

With the digitised outer-wall temperature prescribed, the response of the tube model is controlled almost entirely by the bed-side film coefficient, and the first finding of the study is that the heat-transfer chain printed in Part II of Xu and Froment [6] does not reproduce their own figure. Evaluated as printed, the De Wasch-Froment wall coefficient with the Yagi-Kunii dynamic terms gives $\alpha_i \approx 3100\text{--}3700 \text{ W m}^{-2} \text{ K}^{-1}$ at the particle Reynolds number of the case ($Re_p \approx 4000$), and the Leva-Grummer correlation gives $940\text{--}1580 \text{ W m}^{-2} \text{ K}^{-1}$; the value back-calculated from the digitised wall and gas temperatures and the wall conduction is $348 \text{ W m}^{-2} \text{ K}^{-1}$ (interquartile range $313\text{--}363$). Used as printed, the correlations overpredict the gas temperature by $70\text{--}115 \text{ K}$ and the conversion by $0.14\text{--}0.22$, an RMSE of 75.4 K and 0.160 against 14.0 K and 0.039 for the adopted configuration (Table S4). Both dynamic terms are linear in $Re Pr$ and were established for air an order of magnitude below industrial Reynolds numbers [31, 32], so the extrapolation drives the effective radial conductivity of the bed to some $50 \text{ W m}^{-1} \text{ K}^{-1}$, which no measurement supports. A displaced decimal in the printed effective-conductivity coefficient (0.014 instead of 0.14) would bring the chain to within 30% of the back-calculated value, but this is a hypothesis rather than a correction.




Figure 3: Coupled furnace–tube model after calibration (α_{top} fixed) on the plant cases of Latham [17]: axial profiles of furnace-gas, outer-wall, inner-wall and process-gas temperatures for Plants A, B and C2 with the measured wall temperatures at the two peep-hole elevations, the outlet temperature and the flue-gas outlet temperature (error bars $\pm 1.96\sigma$), and the parity of predicted against measured temperatures (dotted lines ± 10 K).

The Xu–Froment case is therefore used as a verification of the tube model with a matched film coefficient rather than as a validation of the heat-transfer correlation. With $\alpha_i = 348 \text{ W m}^{-2} \text{ K}^{-1}$ and Latham’s inlet effectiveness profile, the model reproduces the digitised profiles with an RMSE of 14.0 K in gas temperature, 16.7 K in inner-wall temperature and 0.015 in the conversion increment along the tube (Fig. 2, Table S4). The residual offset of about 0.04 in absolute conversion originates in the first half metre and reflects an ambiguity in the digitised inlet that the increment metric removes (Section S4). This verifies the kinetics, the thermodynamics, the energy balance and the wall-conduction chain. The pressure RMSE rises from 0.04 bar with the Leva–Grummer run to 0.25 bar with the matched coefficient, for the reason given in Section S5.

The bed-side coefficient is validated instead on the plant data, where the furnace model supplies the outer-wall condition and f_{htg} is a free parameter. Fig. 3, Table 1 and Table S5 summarise the calibration. The adopted fit (α_{top} fixed) gives $F_{\text{gt}} = 0.323$ [0.309, 0.338], $L_q = 0.456$ [0.434, 0.478] and $f_{\text{htg}} = 1.83$ [1.617, 2.052] as 95 % intervals, a χ^2 of 34.9 on 24 degrees of freedom (reduced χ^2 1.46), and reproduces the outlet temperatures of all three cases within 4 K, the lower wall temperatures within 1.1 K and the upper wall temperatures within 12.2 K. The four-parameter fit lowers χ^2 to 30.5 but is degenerate in L_q and α_{top} (Table S5); fixing α_{top} raises χ^2 by 14 %, and the three-parameter fit is adopted for all subsequent work because its parameters are identifiable and its cross-validation folds return nearly the same values (F_{gt} 0.321–0.326, L_q 0.449–0.465, f_{htg} 1.79–1.85). The heat-transfer factor of 1.83 is within 9 % of the 1.68 obtained by Latham et al. [2] with a different furnace model on the same plant, and corresponds to a bed-side coefficient of about $1400 \text{ W m}^{-2} \text{ K}^{-1}$ at mid-height of the Plant A tube, i.e. between the two literature correlations and four times the value that reproduces the Xu–Froment figure; the range $300\text{--}3000 \text{ W m}^{-2} \text{ K}^{-1}$ is the epistemic span carried into the uncertainty analysis through f_{htg} .

Leave-one-case-out cross-validation (last two columns of Table 1) shows that the held-out upper-wall temperature is predicted within 5–18 K and the lower within 2 K. The upper-wall error is systematic: Plant C2, at 71 % rate, is

Table 1: Plant cases of Latham [17] (Appendix H, absolute temperatures) and calibration results: measured inlet and outlet conditions, model errors (predicted minus measured) of the all-cases fit with α_{top} fixed, and leave-one-case-out held-out errors. Plant C1 is excluded from the fit (inconsistent feed flow). Per-tube feed; steam-to-carbon ratio on total hydrocarbon carbon; upper and lower wall temperatures at 3.66 and 8.53 m of 12.5 m.

Case	Rate (%)	T_{in} (K)	p_{in} (bar)	Feed (kmol h^{-1})	S/C	T_{out} (K)	$T_{\text{w,u}}$ (K)	$T_{\text{w,l}}$ (K)	ΔT_{out}	$\Delta T_{\text{w,u}}$	$\Delta T_{\text{w,l}}$	CV $\Delta T_{\text{w,u}}$	CV $\Delta T_{\text{w,l}}$
Plant A	99	884.5	30.06	23.5	2.89	1105.5	1095.9	1132.0	+3.6	-8.4	+0.2	-12.8	+0.7
Plant B	92	887.0	29.44	21.1	2.79	1107.2	1082.3	1130.8	+2.2	-3.6	-1.1	-5.4	-1.7
Plant C1	96	890.9	30.27	28.7	2.85	1122.0	1111.2	1151.5	-	-	-	-	-
Plant C2	71	879.9	28.09	15.9	2.86	1108.9	1075.7	1126.2	+3.9	+12.2	-0.5	+18.2	-0.1




Figure 4: Outer-wall and mid-wall temperature profiles of the calibrated Plant A tube and the corresponding local rupture time from the placeholder Larson–Miller curve at the local hoop stress; the hot spot at $z/L \approx 0.43$ governs the life.

predicted too hot at the upper elevation and Plant A too cold. The upper peep hole lies in the flame zone, where a single parabolic release profile with a load-independent flame length cannot follow the change in flame shape with firing rate. This is the main structural limitation of the furnace model; the residuals are nevertheless within the 15–22 K standard deviation of the wall temperature across the tube population at that elevation [2]. The consequence for the conclusions is examined in Section 3.5.

Fig. 4 shows the life profile of the calibrated Plant A tube. The hot spot lies at about 0.43 of the heated length, where the outer wall reaches 1147 K at a hoop stress of 12.9 MPa; the rupture time from the placeholder minimum curve is 10^8 h at this point, and evaluating the life at the tube-average temperature instead of at the hot spot overstates it by a factor of about 110, which reproduces the contrast reported by Yeh [24]. The absolute value is not credible as a tube life and illustrates why the placeholder must be replaced: the manufacturer’s curve carries no data below about 20 MPa, and the low hoop stress of a thick tube at 30 bar places the operating point in the extrapolated region of the master curve. All results below are therefore expressed as ratios of consumption rates, which depend on the slope of the master curve at the operating point rather than on its absolute position.

3.2. RQ1: what governs hot-spot temperature and life

The closed-loop maps (Fig. 5) show the structure of the operating space when the outlet temperature is held by the firing controller. Across the whole grid the hot spot spans only 1128–1146 K. It rises mildly with load: the duty




Figure 5: Closed-loop operating maps in load and steam-to-carbon ratio at constant outlet temperature (1105.55 K): (a) hot-spot temperature, (b) $\log_{10}$ of the life consumption relative to the base case and (c) net hydrogen per tube with contours of the firing factor required. The star marks the calibrated Plant A base.




Figure 6: First-order (S_1) and total-effect (S_T) Sobol indices with 95 % confidence intervals for (a) the hot-spot temperature and (b) $\log_{10} t_r$ at the hot spot, from 4096 physics evaluations of a Saltelli design over the seven operating inputs.

per tube is proportional to the load, and the bed-side coefficient rises almost proportionally with it ($\alpha_i \propto Re^{0.9}$ in the Leva–Grummer correlation), so the film temperature difference is nearly load-independent and only the conduction drop across the wall grows with the flux. It falls with steam-to-carbon ratio because, at the same outlet temperature, the larger process-gas flow raises the bed-side coefficient, the additional steam increases the gas heat capacity, and the equilibrium shift lets the same methane slip be reached with a slightly lower gas temperature, all of which reduce the wall-to-gas temperature difference despite the higher firing needed. The life map (panel b) is the image of the temperature map through the Larson–Miller relation and spans about half a decade across the space. Hydrogen production (panel c) increases with both load and steam-to-carbon ratio, the latter through the shift of equilibrium, at the cost of a higher firing factor.

The Sobol analysis over all seven inputs (Fig. 6, Table 2) ranks the drivers of the hot-spot temperature: the specific firing factor ($S_T = 0.35$), the steam-to-carbon ratio (0.28) and the excess air (0.19) together explain 82 % of the variance; inlet temperature, catalyst activity and load contribute below 0.1 each, and inlet pressure is negligible. Excess air matters because at fixed fuel it lowers the adiabatic flame temperature and raises the flue-gas flow, both of which move heat release down the tube. The ranking for $\log_{10} t_r$ is the same except that the inlet pressure enters with $S_T = 0.16$ through the hoop stress, which is invisible in the temperature. First-order and total-effect indices agree within their confidence intervals for every input, so interactions are weak over the operating space and the influence of each variable can be read from one-at-a-time maps. The methane slip is governed by the firing factor ($S_T = 0.54$) and the load, as expected for an approach-to-equilibrium quantity.

The surrogates that make the following campaigns possible are summarised in Fig. 7 and Table S3. The Gaussian process is selected for seven of the eight targets, and the physics features improve the hot-spot temperature, outlet temperature and slip. The test RMSE is 0.11 K for the hot-spot temperature and 0.0057 decades for $\log_{10} t_r$, two

Table 2: Sobol first-order (S_1) and total-effect (S_T) indices with 95 % confidence half-widths from a Saltelli design ($N = 256$, 4096 physics runs) over the operating space of Table S1.

Input	$S_1 (T_{\text{wo,max}})$	$S_T (T_{\text{wo,max}})$	$S_1 (\log_{10} t_r)$	$S_T (\log_{10} t_r)$	$S_1 (\text{CH}_4 \text{ slip})$	$S_T (\text{CH}_4 \text{ slip})$
steam-to-carbon	0.29 ± 0.09	0.28 ± 0.05	0.25 ± 0.08	0.25 ± 0.04	0.04 ± 0.03	0.04 ± 0.01
load	0.03 ± 0.03	0.03 ± 0.01	0.03 ± 0.03	0.03 ± 0.01	0.19 ± 0.07	0.19 ± 0.03
firing	0.34 ± 0.09	0.35 ± 0.06	0.29 ± 0.09	0.29 ± 0.05	0.54 ± 0.12	0.54 ± 0.10
inlet T	0.06 ± 0.05	0.06 ± 0.01	0.05 ± 0.04	0.05 ± 0.01	0.08 ± 0.06	0.09 ± 0.02
inlet p	0.02 ± 0.02	0.02 ± 0.00	0.16 ± 0.06	0.16 ± 0.03	0.03 ± 0.03	0.03 ± 0.01
activity	0.07 ± 0.04	0.08 ± 0.02	0.06 ± 0.05	0.06 ± 0.01	0.00 ± 0.01	0.00 ± 0.00
excess air	0.19 ± 0.06	0.19 ± 0.03	0.16 ± 0.07	0.16 ± 0.03	0.11 ± 0.04	0.11 ± 0.02




Figure 7: Parity of the selected surrogates against the physics model on the independent Saltelli test set for (a) the hot-spot temperature, (b) dry methane slip and (c) $\log_{10} t_r$ at the hot spot, with nominal 90 % CV+ conformal intervals.

orders of magnitude below the calibration residuals, at a speed-up of about 4700 relative to the physics model. The position of the hot spot is not learnable ($R^2 = 0.64$) because it is bimodal, an interior maximum near 0.43 of the length for most inputs and an outlet plateau for others, and it is not needed for the life calculation. The CV+ intervals over-cover: at a nominal level of 0.90 the empirical coverage is 0.94–0.97 for all targets. Over-coverage is expected for CV+, whose guarantee is one-sided, and is conservative; the interval widths are nevertheless small enough (0.25 K for the hot-spot temperature) that surrogate error is immaterial to the conclusions.

3.3. RQ2: life cost of flexible operation and ageing

Table 3 and Fig. 8 answer RQ2 for one year of operation at the calibrated base. Under the outlet-temperature controller, load-following costs no tube life per unit of hydrogen: the daily profile (S2) and the renewable-following profile (S3) consume 0.991 and 0.983 of the S1 life per kmol H_2 , i.e. slightly *less*, because the hot-spot temperature at constant outlet temperature falls with load (Fig. 5a) and the hours at reduced load more than compensate the hours at overload. Hydrogen output falls by 10 % and 15 % respectively, so the absolute annual damage falls in proportion. A steady-state assessment at the annual-mean inputs reproduces the hourly accumulation within 2 % for every scenario without discrete events (ratio 0.982 for S3), which means that the non-linearity of the Larson–Miller relation is not strong enough over the load range of this controller to make time-resolved simulation necessary for smooth load variations. The conclusion is conditional on the controller: without outlet-temperature (or slip) control, a load reduction at unchanged firing raises the hot spot sharply, and the open-loop design shows life-consumption ratios of up to two orders of magnitude above the base when high firing coincides with a hot inlet; the operating discipline of controlling on the outlet is itself the main protection of the tubes.

Discrete over-firing events are different in kind. During a mild event (+10 % firing) the damage rate is 2.4 times the base rate, during a severe event (+25 %) it is 6.9 times, and a 40 K hot band raises it by a factor of 11.1; but because the events are short, their share of the annual damage is 0.4–2.7 % and the annual life cost per kmol rises by at most 3 %. The annual-mean evaluation underestimates the event scenarios (diamonds below bars in Fig. 8), as it must for a convex damage law, but the error is small at these event frequencies. The steam-to-carbon ratio has a larger




Figure 8: Annual life consumption per kmol of hydrogen relative to the steady base S1 for the RQ2 scenarios, from hourly accumulation (bars) and from a single evaluation at the annual-mean inputs (diamonds).

Table 3: RQ2 scenario summary for one year (8760 h) of the Plant A twin at 21.6 % excess air: hydrogen per tube, Robinson damage D , life consumption per kmol H_2 relative to S1, ratio of the annual-mean-condition damage estimate to the hourly-accumulated damage, mean and maximum hot-spot temperature and mean firing factor. Life uses the placeholder curve.

Scenario	H_2 (10^3 kmol)	D (-)	Rel. life cost	Avg.-cond. ratio	$\bar{T}_{wo,max}$ (K)	max $T_{wo,max}$ (K)	Firing
S1 steady base	128.5	9.08×10^{-5}	1.000	1.000	1143	1143	0.978
S2 daily load-following	115.7	8.10×10^{-5}	0.991	0.982	1141	1143	0.944
S3 renewable-following	109.7	7.61×10^{-5}	0.983	0.982	1140	1146	0.927
S4a mild over-firing (48 h)	128.6	9.14×10^{-5}	1.007	0.993	1144	1157	0.979
S4b severe over-firing (6 h)	128.5	9.11×10^{-5}	1.004	0.996	1143	1179	0.978
S4c hot band +40 K (24 h)	128.5	9.33×10^{-5}	1.028	0.980	1144	1183	0.978
S6 steam-to-carbon 2.5	119.9	9.83×10^{-5}	1.161	1.000	1144	1144	0.897
S6 steam-to-carbon 3.5	139.8	8.23×10^{-5}	0.834	1.000	1142	1142	1.092
S5 ageing, hold slip, year 1	128.5	6.49×10^{-5}	0.715	0.997	1138	1139	0.974
S5 ageing, hold slip, year 4	128.5	2.05×10^{-4}	2.255	0.949	1156	1164	0.989
S5 ageing, hold T_{out} , year 1	129.0	7.00×10^{-5}	0.768	0.998	1139	1140	0.982
S5 ageing, hold T_{out} , year 4	127.4	1.69×10^{-4}	1.875	0.967	1153	1159	0.969

steady effect: operating at 2.5 costs 1.16 and at 3.5 costs 0.83 of the S1 life per kmol H_2 , because the higher ratio both cools the hot spot and raises the hydrogen yield.

Catalyst ageing dominates every operational effect. With the activity falling from 0.30 to 0.10 over four years, a range that brackets the mid-campaign value of 0.20 fitted to the plant and is consistent with the rapid initial deactivation followed by a slow decline observed for fresh catalyst [5] and with the four-year catalyst life of the reference plant [2], and with the controller holding the methane slip, the life cost per kmol H_2 rises from 0.71 in the first year (fresh catalyst) to 2.25 in the fourth as the hot spot climbs from 1138 to 1164 K; holding the outlet temperature instead gives 1.88 in the fourth year because the slip is allowed to rise. Over the campaign the life consumed in year 4 is three times that in year 1, and the question of when to change the catalyst is therefore also a tube-life question.

3.4. RQ3: Pareto-optimal regimes and the steam credit

The Pareto set (Fig. 9, Table 4) contains 200 members, of which 199 remain feasible under the physics model; the maximum discrepancy between surrogate and physics on the set is 0.03 in relative life and below 0.1 MJ kmol^{-1} in fuel. The calibrated base is dominated by 37 members. The knee of the front raises hydrogen by +16 % while cutting fuel per kmol by -18 % and life consumption per kmol to 0.28 of the base; it does so by running at 110 % load with the minimum excess air, the maximum inlet temperature and a lower firing factor. At constant production (iso- H_2 regime) the life cost can be reduced to 0.083 of the base with -16 % fuel by reducing the firing and letting the hot spot fall from 1147 to 1109 K. The maximum-hydrogen corner (+30 %) costs 4.3 times the base life per kmol and runs close to the wall-temperature limit.




Figure 9: Physics-verified NSGA-II Pareto set in (a) firebox fuel per kmol H_2 against life consumption per kmol H_2 , (b) net hydrogen per tube against life consumption and (c) hydrogen against fuel; open circles are members infeasible under the physics model, and the marked regimes are listed in Table 4.

Table 4: RQ3 representative regimes (physics-verified) from the NSGA-II Pareto set (200×300 generations; 199 of 200 members feasible under the physics model). Fuel is firebox fuel per kmol H_2 ; life is the consumption rate per kmol H_2 relative to the base (placeholder curve).

Regime	Load	S/C	T_{in} (K)	Excess air (%)	Firing	H_2 (kmol h^{-1})	Fuel (MJ kmol $^{-1}$ H_2)	Rel. life cost	$T_{wo,max}$ (K)	CH_4 slip (%)
base	1.00	2.89	885	21.6	1.000	14.82	158.4	1.000	1147	7.81
min life, iso- H_2	0.96	4.00	895	17.5	0.873	14.72	133.7	0.083	1109	7.77
knee	1.10	4.00	900	5.0	0.871	17.22	130.5	0.280	1130	7.15
max H_2	1.10	4.00	900	5.0	1.150	19.26	154.1	4.329	1177	3.89
min fuel overall	0.60	4.00	900	5.0	0.850	9.87	121.3	0.472	1128	5.72
min life overall	0.60	4.00	858	24.9	0.850	9.18	130.4	0.046	1091	7.81

Every member of the front sits at the upper bound of the steam-to-carbon ratio (3.96–4.00). This is an artefact of the objective: with firebox fuel alone, additional steam is free, and the equilibrium shift and the cooling of the hot spot both favour maximum steam. The steam-credit sweep (Fig. 10, Table S6) shows how the answer changes when steam raising and feed preheat are charged. With steam fully charged ($\alpha = 0$) the minimum-fuel regime moves to a steam-to-carbon ratio of 2.66 and the knee to 3.3, while the minimum-life regime stays at 4.0 but now costs +3.2 % more total fuel than the base for a life cost of 0.059. For $\alpha \geq 0.5$, i.e. when at least half of the marginal steam duty is recovered from the flue gas or credited as export steam, the ratio of 4.0 is optimal for all three regimes. In a hydrogen plant the process steam is raised in the waste-heat boiler on the reformer effluent and in the flue-gas convection section, and the surplus is exported, so the marginal steam for a higher ratio mostly displaces export steam rather than requiring extra fuel; this corresponds to α between about 0.5 and 1 depending on the value the site attaches to export steam, and places most plants in the region where tube life and fuel economy pull in the same direction. The plant’s own steam balance, not the reformer, therefore decides whether running steam-rich is a life-saving measure or a fuel penalty, and the twin makes the dependence explicit.

3.5. RQ4: uncertainty and robustness

Table 5 and Fig. 11 give the propagated uncertainty. At the base point the 90 % interval of the hot-spot temperature is 42 K wide and that of $\log_{10} t_r$ spans 7.04–8.73, i.e. a factor of 50 in absolute rupture time. The paired ratios are much tighter: the knee regime costs 0.20–0.64 of the base life per kmol and the fourth ageing year 1.8–9.0 times, so the ranking of regimes survives an uncertainty that makes the absolute life meaningless. The nominal result lies inside the 90 % interval at every point, and the end-of-campaign hot spot exceeds the 1193 K limit in about 8 % of the samples. The variance decomposition (Fig. 12, Table S7) attributes 62 % of the variance of $\log_{10} t_r$ to the creep group, dominated by the heat-to-heat scatter and the Larson–Miller constant, 23 % to the kinetics, 9 % to the wall-temperature measurement and only 5 % to the calibrated heat-transfer parameters; for the hot-spot temperature the kinetics group dominates (51 %), followed by the measurement offset and the wall thickness. The one-at-a-time shares sum to within 0.04 of one, and the Sobol total effects on the surrogate confirm the ranking, so the groups act almost additively.

The robustness fractions (Table 6) state which conclusions hold across the joint uncertainty. That the fourth ageing year costs more than twice the S1 life per kmol holds in 0.91 of the samples, that the knee regime halves the life cost




Figure 10: Iso-production front as a function of the steam credit α : (a) optimal steam-to-carbon ratio and (b) life consumption per kmol H_2 relative to the base for the minimum-life, knee and minimum-fuel regimes. The full sweep is given in Table S6.




Figure 11: Monte Carlo distributions (2048 samples of all uncertainty groups through the physics model) of (a) the hot-spot temperature, with the 1193 K limit, and (b) $\log_{10} t_r$ at the hot spot, at the S1 base, the RQ3 knee, the iso- H_2 minimum-life regime and the end of the fourth ageing year.

in 0.81, and that steam-to-carbon 3.5 is cheaper in life than 2.5 in 0.84; these three conclusions are robust and, for the knee, become more so under the flame-length variant. The conclusion that flexible operation costs no life per kmol is not robust: it holds in 0.57–0.63 of the samples with the calibrated flame length, close to a coin toss because the nominal effect is within 2 % of unity, and it is reversed in all samples under the bounding variant, in which the hot spot at 70 % load moves up the tube into the flame zone and the life cost per kmol rises to 1.71 times the base. The sign of the flexible-operation effect is therefore determined by a structural feature of the furnace, the response of the heat-release profile to load, that the three plant cases cannot identify.

4. Discussion

4.1. Industrial implications

Three results have direct operational consequences. First, the drivers of tube life are the firing intensity, the steam-to-carbon ratio and the excess air, in that order, and they act almost independently (Section 3.2). Of these, excess air is the least appreciated: at fixed fuel it lowers the hot-spot temperature by shifting heat release down the tube, and the optimiser exploits it in both directions, choosing minimum excess air when hydrogen is the priority and maximum excess air when life is. Its efficiency penalty, the larger flue-gas flow and stack loss, is included in the model through the firing required to hold the outlet temperature, so the optimiser weighs it against life rather than ignoring it. Excess air is also the cheapest variable to move on an operating plant, and its effect can be verified with the existing pyrometer survey.

Second, the trade-off between production, fuel and life has a knee that is far from the current operating point of the calibrated plant. The base is dominated by a substantial part of the front, and the knee regime obtains more

Table 5: RQ4 Monte Carlo (scrambled Sobol design, $N = 2048$, all uncertainty groups) through the physics model: median and 5–95 % interval of the hot-spot temperature, of $\log_{10} t_r$ and of the paired life consumption per kmol H_2 relative to the base.

Operating point	$T_{wo,max}$ (K)	$\log_{10} t_r$ (h)	Rel. life cost
S1 base	1145 [1127, 1169]	7.91 [7.04, 8.73]	1.000 (paired)
RQ3 knee	1132 [1107, 1165]	8.30 [7.33, 9.20]	0.338 [0.201, 0.643]
RQ3 min-life, iso- H_2	1110 [1088, 1138]	8.90 [7.94, 9.80]	0.100 [0.069, 0.159]
S5 year 4, hold slip	1165 [1138, 1203]	7.35 [6.28, 8.30]	3.478 [1.841, 8.974]
S/C 2.5	1146 [1129, 1166]	7.88 [7.04, 8.69]	1.153 [0.916, 1.350]
S/C 3.5	1144 [1127, 1172]	7.94 [7.04, 8.83]	0.846 [0.587, 1.119]
load 0.70	1138 [1123, 1156]	8.08 [7.22, 8.90]	0.966 [0.638, 1.350]
load 0.85	1141 [1124, 1162]	8.00 [7.14, 8.82]	0.951 [0.787, 1.143]




Figure 12: One-group-at-a-time shares of the variance of $\log_{10} t_r$ and of the hot-spot temperature at the S1 base point (numerical values in Table S7).

hydrogen at less fuel per kmol and a fraction of the life cost, mainly by running at higher load with lower specific firing and the highest inlet temperature. Whether these gains are realisable depends on constraints outside the twin (preheat-coil metallurgy at 900 K inlet temperature, pressure-swing-adsorption capacity at 110 % load, downstream steam balance), and the steam-credit sweep shows that the recommended steam-to-carbon ratio moves from 4.0 to about 2.7 depending on how the plant values steam. The twin cannot decide this, but it makes the dependence explicit, which is the useful output.

Third, the life cost of load-following is small compared with the life cost of catalyst ageing, provided the outlet temperature is controlled. The practical message is that a plant that must follow demand should not fear the tube-life consequences of smooth load changes, but should watch the end-of-campaign hot-spot temperature under a methane-slip target, where the twin predicts the fourth year to cost 2.3 times the first per kmol of hydrogen. Under uncertainty this ageing conclusion is the most robust of the study, and the conclusion on load-following the least. The difference between the two comes from the size of the nominal effect (a factor of 2.3 against 2 %), not from a difference in modelling quality, and we recommend that digital-twin studies report robustness fractions alongside nominal results so that small nominal effects are not mistaken for findings.

4.2. The flame-length response to load as a measurement need

The one structural assumption that changes a conclusion is how the heat-release profile of the burners responds to the firing rate. The calibrated model uses a single dimensionless flame length, and the cross-validation residuals

Table 6: Robustness of the RQ2/RQ3 conclusions: fraction of paired Monte Carlo samples supporting each statement, with the calibrated flame length L_q and with the bounding structural variant $L_q \text{ load}^{0.5}$.

Statement	nominal L_q	$L_q \propto \text{load}^{0.5}$
ageing year 4 costs $> 2 \times$ S1 life per kmol	0.906	0.906
knee regime $< 0.5 \times$ base	0.806	0.936
S/C 3.5 cheaper in life than S/C 2.5	0.836	0.836
S2 daily $\leq$ S1	0.568	0.000
S3 renewable $\leq$ S1	0.632	0.001

at the upper peep hole (Section 3.1) already show that this is the weakest element of the model: the partially loaded Plant C2 is predicted too hot at the top and the fully loaded Plant A too cold. When the flame is allowed to shorten with load, the flexible-operation conclusion reverses because the hot spot at reduced load moves into a hotter part of the furnace. The two hypotheses correspond to two ways of turning a furnace down: reducing the fuel flow through all burners, which leaves the length of a turbulent diffusion flame nearly unchanged [35], or taking burner rows out of service, which redistributes the heat release; the three published operating points cannot separate them. What would settle the question is a pyrometer survey at three or four elevations at two or three firing rates on one furnace, which most plants already perform in part during turnaround planning; a digital twin of the kind presented here turns such a survey into an estimate of the flame-length response with its uncertainty, and the robustness fractions would then no longer straddle 0.5. We regard this as the single most valuable measurement for tube-life-aware operation of top-fired reformers.

4.3. Limitations

The model represents one average tube. Tube-to-tube variation in a top-fired furnace, caused by burner imbalance, row position and flow maldistribution, is of the order of 15–22 K (one standard deviation across the tube population at a given elevation in the reference furnace [2]), which by the Larson–Miller relation corresponds to a factor of two to four in rupture time between the hottest and the average tube. The measurement group of the uncertainty analysis (± 10 K) partly covers this, and a per-tube extension would replace the single wall-temperature offset by a distribution over the tube population, but the present results are statements about an average tube.

The life model covers creep only. Thermal fatigue during start-up and shutdown, creep–fatigue interaction and carburisation are not modelled, and the hoop stress is the pressure stress alone. The elastic thermal stress from the radial temperature gradient, $E\alpha\Delta T_{\text{wall}}/[2(1-\nu)]$, is about 45 MPa for the 20–30 K wall drop of the calibrated tube ($E \approx 150$ GPa, $\alpha \approx 18 \times 10^{-6}$ K $^{-1}$ at 900 °C), several times the 12 MPa hoop stress; in steady operation it relaxes by creep and is conventionally excluded from the design stress, but it is the driver of damage during thermal cycling. The scenario results should therefore be read as the creep-life cost of the scenarios, not their total life cost, and the ranking of flexible-operation scenarios could change if the number of thermal cycles were counted.

The rupture curve is a placeholder. The manufacturer’s minimum curve for a proprietary alloy reproduced by Yeh [24] was used because no public curve with scatter information was available in the transcribed form the module requires, and the absolute rupture times it yields at the low hoop stress of this tube are not credible. The uncertainty analysis compensates with a factor-of-two heat-to-heat scatter and a ± 1 range on the Larson–Miller constant, and the creep group is the largest contributor to the variance of $\log_{10} t_r$; replacing the curve by a fit to alloy-specific rupture data with scatter, from manufacturers’ average and minimum design curves or from tabulated rupture points in the open literature [21, 22], is the first step towards any absolute statement. The module is built for this replacement and the ingestion has been tested on synthetic data.

The scenarios are quasi-steady: each hour is a steady state of the twin, and the thermal inertia of the tube, the catalyst and the furnace refractory is neglected. The thermal time constant of the tube wall ($\rho c t^2/\lambda$) is of the order of one minute and that of the refractory lining of the order of a day, so the tube follows hourly load changes essentially instantaneously while the refractory lags them; for load ramps of 0.1 per hour and events of one hour or longer the quasi-steady approximation is reasonable for the tube, but it excludes the transient overshoot of the wall temperature during a fast firing step, which would increase the cost of the over-firing events. Finally, the calibration rests on three operating points of one furnace from a single published source, so the parameter intervals are conditional on that

furnace and the transfer to another reformer requires recalibration on its own survey data; the pipeline is designed so that this is a matter of replacing one CSV file.

5. Conclusions

A physics-informed digital twin of an industrial top-fired steam methane reformer has been built from published kinetics, plant data and creep curves, coupling a one-dimensional tube model, a long-furnace model, a Larson–Miller life module, Gaussian-process surrogates with conformal prediction intervals and Monte Carlo uncertainty propagation in one open, versioned pipeline.

- The tube model reproduces the industrial case of Xu and Froment within 14 K in gas temperature once the film coefficient is matched, and the heat-transfer correlations printed in that paper are found to overpredict the film coefficient by an order of magnitude at industrial Reynolds numbers. The coupled model calibrated on three plant operating points reproduces outlet temperatures within 4 K and wall temperatures within 12 K, with identifiable parameters and a systematic residual at the upper elevation that points to the flame-length response to load.
- Firing intensity, steam-to-carbon ratio and excess air explain 82 % of the variance of the hot-spot temperature with negligible interactions; inlet pressure enters the life only through the hoop stress.
- Under outlet-temperature control, load-following and renewable-following cost no life per kmol of hydrogen, short over-firing events at most 3 %, and the steam-to-carbon ratio ± 17 %; catalyst ageing under a methane-slip target multiplies the life cost by 2.25 in four years.
- The calibrated base is Pareto-dominated: a knee regime gives +16 % hydrogen at –18 % fuel per kmol and 0.28 of the life cost, and at constant production the life cost can be cut to 0.083. The optimal steam-to-carbon ratio depends on the steam credit, moving from 4.0 to 2.7 when steam raising is charged in full.
- The 90 % uncertainty on the absolute rupture time spans a factor of 50, dominated by creep-curve scatter (62 %), while paired ratios between regimes remain tight. The ageing, knee and steam-to-carbon conclusions hold in 81 % to 91 % of the samples; the flexible-operation conclusion does not survive a load-dependent flame length.

The next steps are the ingestion of alloy-specific rupture data with scatter for the HK40 and HP-Nb alloys, a per-tube extension with a measured wall-temperature distribution, and, above all, a multi-elevation pyrometer survey at several firing rates to identify the flame-length response that decides the life cost of flexible operation.

Data and code availability

The model code (`rdt` package), the pipeline entry point (`python -m rdt.pipeline -stage all -version v2`), the transcribed literature data with their provenance annotations, the digitised verification curves, all v2 result files and the scripts that generate every figure, table and number in this manuscript are available at <https://github.com/IzbassarO/reformer-digital-twin> and archived at [TODO: Zenodo DOI]. The literature PDFs are not redistributed. Transcribed data are released under CC BY 4.0 and code under the MIT licence.

CRediT authorship contribution statement

Izbassar Orynbassar: Conceptualisation, Methodology, Software, Validation, Formal analysis, Data curation, Visualisation, Writing – original draft, Writing – review & editing. **Madina Sissenbay:** Conceptualisation (process and materials scope), Validation (chemical-engineering content), Investigation (creep-data transcription), Writing – review & editing.

Declaration of generative AI and AI-assisted technologies in the writing process

During the preparation of this work the authors used Claude (Anthropic) in order to assist with code development, data transcription checks, figure and table generation and drafting of the text. After using this tool, the authors reviewed and edited the content as needed and take full responsibility for the content of the published article.

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper. M.S. is employed in the industrial gas sector; the work uses exclusively published data and no proprietary information from any employer.

Funding

This research did not receive any specific grant from funding agencies in the public, commercial, or not-for-profit sectors.

Acknowledgements

I.O. acknowledges the support of the Bolashak International Scholarship of the Republic of Kazakhstan during the preparation of this work. The authors thank D. A. Latham for the publicly archived thesis appendix that made the plant-scale calibration possible.

References

- [1] International Energy Agency, Global Hydrogen Review 2024, IEA, Paris, 2024.
- [2] D. A. Latham, K. B. McAuley, B. A. Peppley, T. M. Raybold, Mathematical modeling of an industrial steam-methane reformer for on-line deployment, *Fuel Process. Technol.* 92 (2011) 1574–1586. <https://doi.org/10.1016/j.fuproc.2011.04.001>
- [3] A. K. Ray, S. K. Sinha, Y. N. Tiwari, J. Swaminathan, G. Das, S. Chaudhuri, R. Singh, Analysis of failed reformer tubes, *Eng. Fail. Anal.* 10 (2003) 351–362. [https://doi.org/10.1016/S1350-6307\(02\)00029-8](https://doi.org/10.1016/S1350-6307(02)00029-8)
- [4] American Petroleum Institute, API Standard 530: Calculation of Heater-Tube Thickness in Petroleum Refineries, 7th ed., API, Washington, DC, 2015.
- [5] J. Xu, G. F. Froment, Methane steam reforming, methanation and water-gas shift: I. Intrinsic kinetics, *AIChE J.* 35 (1989) 88–96. <https://doi.org/10.1002/aic.690350109>
- [6] J. Xu, G. F. Froment, Methane steam reforming: II. Diffusional limitations and reactor simulation, *AIChE J.* 35 (1989) 97–103. <https://doi.org/10.1002/aic.690350110>
- [7] G. Pantoleontos, E. S. Kikkinides, M. C. Georgiadis, A heterogeneous dynamic model for the simulation and optimisation of the steam methane reforming reactor, *Int. J. Hydrogen Energy* 37 (2012) 16346–16358. <https://doi.org/10.1016/j.ijhydene.2012.02.125>
- [8] P. M. Plehiers, G. F. Froment, Coupled simulation of heat transfer and reaction in a steam reforming furnace, *Chem. Eng. Technol.* 12 (1989) 20–26. <https://doi.org/10.1002/ceat.270120105>
- [9] A. Zamaniyan, H. Ebrahimi, J. S. S. Mohammadzadeh, A unified model for top fired methane steam reformers using three-dimensional zonal analysis, *Chem. Eng. Process.* 47 (2008) 946–956. <https://doi.org/10.1016/j.cep.2007.03.005>
- [10] A. Kumar, M. Baldea, T. F. Edgar, O. A. Ezekoye, Smart manufacturing approach for efficient operation of industrial steam-methane reformers, *Ind. Eng. Chem. Res.* 54 (2015) 4360–4370. <https://doi.org/10.1021/ie504087z>
- [11] A. Jafarizadeh, M. Panjepour, M. D. Emami, Advanced modelling and optimization of steam methane reforming: from CFD simulation to machine learning-driven optimization, *Int. J. Hydrogen Energy* 96 (2024) 1262–1280. <https://doi.org/10.1016/j.ijhydene.2024.11.352>

- [12] L. E. Yerimah, M. Mehana, R. P. Singh, P. Sharan, Uncertainty-aware hydrogen production optimization from a steam methane reformer with a water gas shift reactor and pressure swing adsorption, *Ind. Eng. Chem. Res.* 64 (2025) 18298–18314. <https://doi.org/10.1021/acs.iecr.5c01681>
- [13] S. R. Nabavi, Z. Guo, Z. Wang, Machine learning-assisted surrogate modeling with multi-objective optimization and decision-making of a steam methane reforming reactor, *arXiv:2507.07641* (2025). <https://arxiv.org/abs/2507.07641>
- [14] National Institute for Materials Science, Data sheets on the elevated-temperature properties of centrifugally cast 25Cr-20Ni-0.4C steel tubes for use in reformer furnaces (SCH 22-CF), NIMS Creep Data Sheet No. 16B, Tsukuba, Japan. <https://doi.org/10.11503/nims.1020>
- [15] National Institute for Materials Science, Data sheets on the elevated-temperature properties of centrifugally cast tubes and cast block of 25Cr-35Ni-0.4C steel for reformer furnaces (SCH 24), NIMS Creep Data Sheet No. 38A, Tsukuba, Japan. <https://doi.org/10.11503/nims.1042>
- [16] National Institute for Materials Science, MatNavi Service Terms of Use, NIMS, Tsukuba, Japan. https://mits.nims.go.jp/agreement/MatNavi_agreement_en.pdf
- [17] D. A. Latham, Mathematical modelling of an industrial steam methane reformer, M.Sc.E. thesis, Department of Chemical Engineering, Queen's University, Kingston, Ontario, Canada, 2008. <http://hdl.handle.net/1974/1650>
- [18] L. Lao, A. Aguirre, A. Tran, Z. Wu, H. Durand, P. D. Christofides, CFD modeling and control of a steam methane reforming reactor, *Chem. Eng. Sci.* 148 (2016) 78–92. <https://doi.org/10.1016/j.ces.2016.03.038>
- [19] A. Tran, A. Aguirre, M. Crose, H. Durand, P. D. Christofides, Temperature balancing in steam methane reforming furnace via an integrated CFD/data-based optimization approach, *Comput. Chem. Eng.* 104 (2017) 185–200. <https://doi.org/10.1016/j.compchemeng.2017.04.013>
- [20] A. Kumar, M. Baldea, T. F. Edgar, A physics-based model for industrial steam-methane reformer optimization with non-uniform temperature field, *Comput. Chem. Eng.* 105 (2017) 224–236. <https://doi.org/10.1016/j.compchemeng.2017.01.002>
- [21] M. Whittaker, B. Wilshire, J. Brear, Creep fracture of the centrifugally-cast superaustenitic steels, HK40 and HP40, *Mater. Sci. Eng. A* 580 (2013) 391–396. <https://doi.org/10.1016/j.msea.2013.05.041>
- [22] G. N. Haidemenopoulos, A. D. Zervaki, H. Kamoutsi, K. Polychronopoulou, Creep rupture in HP-Nb refractory steel tubes due to short-term overheating, *Eur. J. Mater.* 1 (2021) 1–22. <https://doi.org/10.1080/26889277.2021.1994841>
- [23] K. Guguloth, N. Roy, Creep life estimation of reformer alloy using θ -projection method – a neuro fuzzy approach, *Int. J. Press. Vessels Piping* 203 (2023) 104938. <https://doi.org/10.1016/j.ijpvp.2023.104938>
- [24] C.-L. Yeh, Effect of burner operation on the catalyst tube lifetime of a steam methane reformer: a numerical study, *Appl. Sci.* 11 (2021) 231. <https://doi.org/10.3390/app11010231>
- [25] F. R. Larson, J. Miller, A time-temperature relationship for rupture and creep stresses, *Trans. ASME* 74 (1952) 765–775.
- [26] E. L. Robinson, Effect of temperature variation on the long-time rupture strength of steels, *Trans. ASME* 74 (1952) 777–781.
- [27] D. G. Goodwin, H. K. Moffat, I. Schoegl, R. L. Speth, B. W. Weber, Cantera: an object-oriented software toolkit for chemical kinetics, thermodynamics, and transport processes, version 3.1, 2024. <https://doi.org/10.5281/zenodo.14455267>
- [28] S. Ergun, Fluid flow through packed columns, *Chem. Eng. Prog.* 48 (1952) 89–94.
- [29] M. H. Wesenberg, H. F. Svendsen, Mass and heat transfer limitations in a heterogeneous model of a gas-heated steam reformer, *Ind. Eng. Chem. Res.* 46 (2007) 667–676.
- [30] M. Leva, M. Grummer, Heat transfer to gases through packed tubes: effect of particle characteristics, *Ind. Eng. Chem.* 40 (1948) 415–419. <https://doi.org/10.1021/ie50459a012>
- [31] A. P. De Wasch, G. F. Froment, Heat transfer in packed beds, *Chem. Eng. Sci.* 27 (1972) 567–576. [https://doi.org/10.1016/0009-2509\(72\)87012-X](https://doi.org/10.1016/0009-2509(72)87012-X)
- [32] S. Yagi, D. Kunii, Studies on heat transfer near wall surface in packed beds, *AIChE J.* 6 (1960) 97–104. <https://doi.org/10.1002/aic.690060119>

- [33] H. C. Hottel, A. F. Sarofim, Radiative Transfer, McGraw-Hill, New York, 1967.
- [34] F. C. Roesler, Theory of radiative heat transfer in co-current tube furnaces, Chem. Eng. Sci. 22 (1967) 1325–1336. [https://doi.org/10.1016/0009-2509\(67\)80023-X](https://doi.org/10.1016/0009-2509(67)80023-X)
- [35] W. R. Hawthorne, D. S. Weddell, H. C. Hottel, Mixing and combustion in turbulent gas jets, Third Symposium on Combustion and Flame and Explosion Phenomena 3 (1949) 266–288.
- [36] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, É. Duchesnay, Scikit-learn: machine learning in Python, J. Mach. Learn. Res. 12 (2011) 2825–2830.
- [37] R. F. Barber, E. J. Candès, A. Ramdas, R. J. Tibshirani, Predictive inference with the jackknife+, Ann. Stat. 49 (2021) 486–507. <https://doi.org/10.1214/20-AOS1965>
- [38] V. Taquet, V. Blot, T. Morzadec, L. Lacombe, N. Brunel, MAPIE: an open-source library for distribution-free uncertainty quantification, arXiv:2207.12274 (2022). <https://arxiv.org/abs/2207.12274>
- [39] A. Saltelli, P. Annoni, I. Azzini, F. Campolongo, M. Ratto, S. Tarantola, Variance based sensitivity analysis of model output. Design and estimator for the total sensitivity index, Comput. Phys. Commun. 181 (2010) 259–270. <https://doi.org/10.1016/j.cpc.2009.09.018>
- [40] J. Herman, W. Usher, SALib: an open-source Python library for sensitivity analysis, J. Open Source Softw. 2 (2017) 97. <https://doi.org/10.21105/joss.00097>
- [41] K. Deb, A. Pratap, S. Agarwal, T. Meyarivan, A fast and elitist multiobjective genetic algorithm: NSGA-II, IEEE Trans. Evol. Comput. 6 (2002) 182–197. <https://doi.org/10.1109/4235.996017>
- [42] J. Blank, K. Deb, pymoo: multi-objective optimization in Python, IEEE Access 8 (2020) 89497–89509. <https://doi.org/10.1109/ACCESS.2020.2990567>